


# Time Series



DATA SCIENCE BOOTCAMP

Previously, on Linear Regression...



$$y_{\beta}(x) = \beta_0 + \beta_1 x + \varepsilon$$

$$\beta_0 = 94.68 \text{ million}$$

$$\beta_1 = 0.1$$

We expect box office returns for a movie to be  
\$95 Million  
+ 10% of its budget  
+ luck factor  
(what we can't predict)  
(roll the Gaussian dice for about +- \$100 Million)

$$y_{\beta}(x) = \beta_0 + \beta_1 x + \varepsilon$$

$$\begin{aligned}\beta_0 &= 94.68 \text{million} \\ \beta_1 &= 0.1\end{aligned}$$

or something like...

We expect box office returns for a movie to be

**\$20 Million**

+ 25% of its budget

+\$60 Million if it is an action movie in summer

+\$8 Million per 10% rottentomatoes score over 60%

-\$80 Million if directed by Uwe Boll

+ luck factor

(roll the dice for +- \$70 Million)

$$y_{\beta}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_5 x_5 + \varepsilon$$

What we know  
(features)

What we predict  
(target)

Budget

Action in Summer?

RottenTomatoes Score - 60

Director Uwe Boll?



Total Gross

What we know  
(features)

What we predict  
(target)

Budget

Action in Summer?

RottenTomatoes Score - 60

Director Uwe Boll?



Total Gross

+ luck factor  
means this is a  
**stochastic** process

# What if we wanted to predict weekly returns?



## Mad Max: Fury Road

Domestic Total Gross: **\$153,636,354**  
Domestic Lifetime Gross: **\$154,058,340**

Distributor: **Warner Bros.**

Release Date: **May 15, 2015**

Genre: **Sci-Fi Action**

Runtime: **2 hrs. 0 min.**

MPAA Rating: **R**

Production Budget: **\$150 million**



[Get local showtimes at IMDb](#)

[Summary](#)

[Daily](#)

[Weekend](#)

**[Weekly](#)**

[Releases](#)

[Foreign](#)


[Similar Movies](#)

### 2015

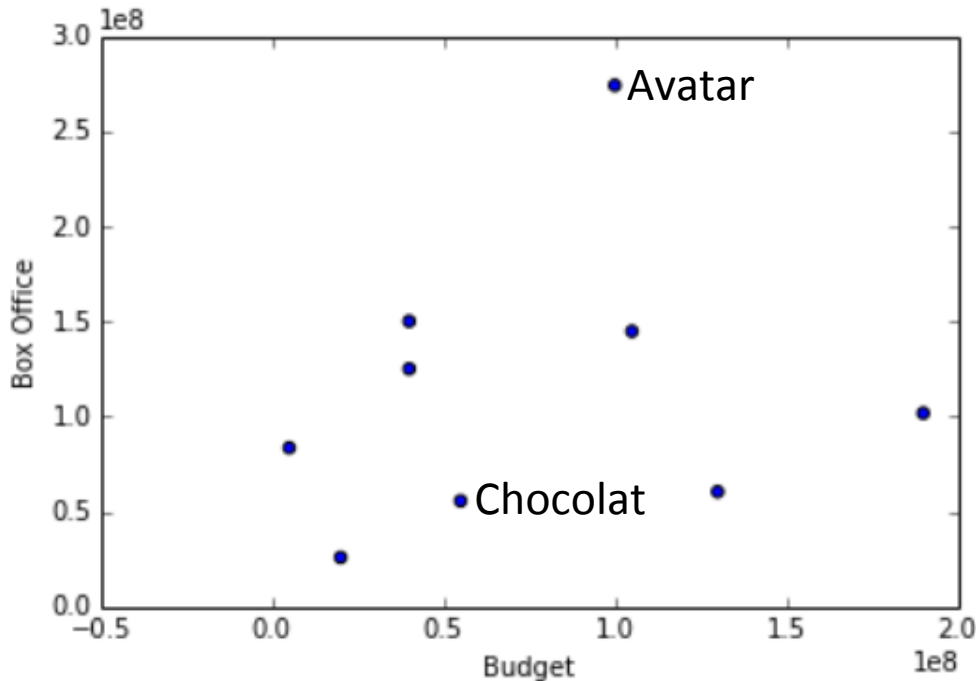
Date (click to view chart)	Rank	Weekly Gross	% Change	Theaters / Change		Avg.	Gross-to-Date	Week #
May 15–21	2	\$63,440,279	-	3,702	-	\$17,137	\$63,440,279	1
May 22–28	3	\$38,849,255	-38.8%	3,722	+20	\$10,438	\$102,299,534	2
May 29–Jun 4	3	\$20,544,731	-47.1%	3,255	-467	\$6,312	\$122,834,265	3
Jun 5–11	5	\$11,643,562	-43.3%	2,720	-535	\$4,281	\$134,477,827	4
Jun 12–18	6	\$6,309,002	-45.8%	2,234	-486	\$2,824	\$140,786,829	5
Jun 19–25	8	\$4,555,993	-27.8%	1,424	-810	\$3,199	\$145,342,822	6
Jun 26–Jul 2	11	\$2,648,047	-41.9%	561	-863	\$4,720	\$147,990,879	7
Jul 3–9	14	\$1,615,110	-37.9%	561	0	\$2,880	\$149,605,989	8



# What if we wanted to predict weekly returns?




In the previous case,  
each point is independent.



Each point is an  
independent 'firing' of  
the stochastic process.

How much Avatar makes does not  
limit or influence by itself how much  
Chocolat makes

In a time series,  
points are **not** independent.



Points carry information  
about other points

Clearly, if the movie has made \$600K on week 10, and \$470K on week 11, you know that the movie will not make \$2M on week 12.

But the approach does not need to be different from any other regression

What we know  
(features)

What we predict  
(target)

Values that carry  
information about  
the target



Target  
Value

# But the approach does not need to be different from any other regression

What we know  
(features)

What we predict  
(target)

Gross of  
Week 10



Gross of  
Week 11

Since past points carry information about the future points, use these past points as features

**Target**



**Feature**



	<b>weekly_gross</b>	<b>one_prev_weeks_gross</b>
<b>0</b>	63440279	NaN
<b>1</b>	38849255	63440279
<b>2</b>	20544731	38849255
<b>3</b>	11643562	20544731
<b>4</b>	6309002	11643562
<b>5</b>	4555993	6309002
<b>6</b>	2648047	4555993
<b>7</b>	1645168	2648047
<b>8</b>	966275	1645168
<b>9</b>	601794	966275
<b>10</b>	663222	601794

## Auto-Regressive Models

Regression, where one or more previously observed target values is used as (a) feature(s)

Target



Feature



	weekly_gross	one_prev_weeks_gross
0	63440279	NaN
1	38849255	63440279
2	20544731	38849255
3	11643562	20544731
4	6309002	11643562
5	4555993	6309002
6	2648047	4555993
7	1645168	2648047
8	966275	1645168
9	601794	966275
10	663222	601794

AR-1

*AutoRegressive model with lag 1*

Just one previous point is among  
the features

# AR-1

What we know  
(features)

What we predict  
(target)

Gross of  
previous week



Gross of  
this week



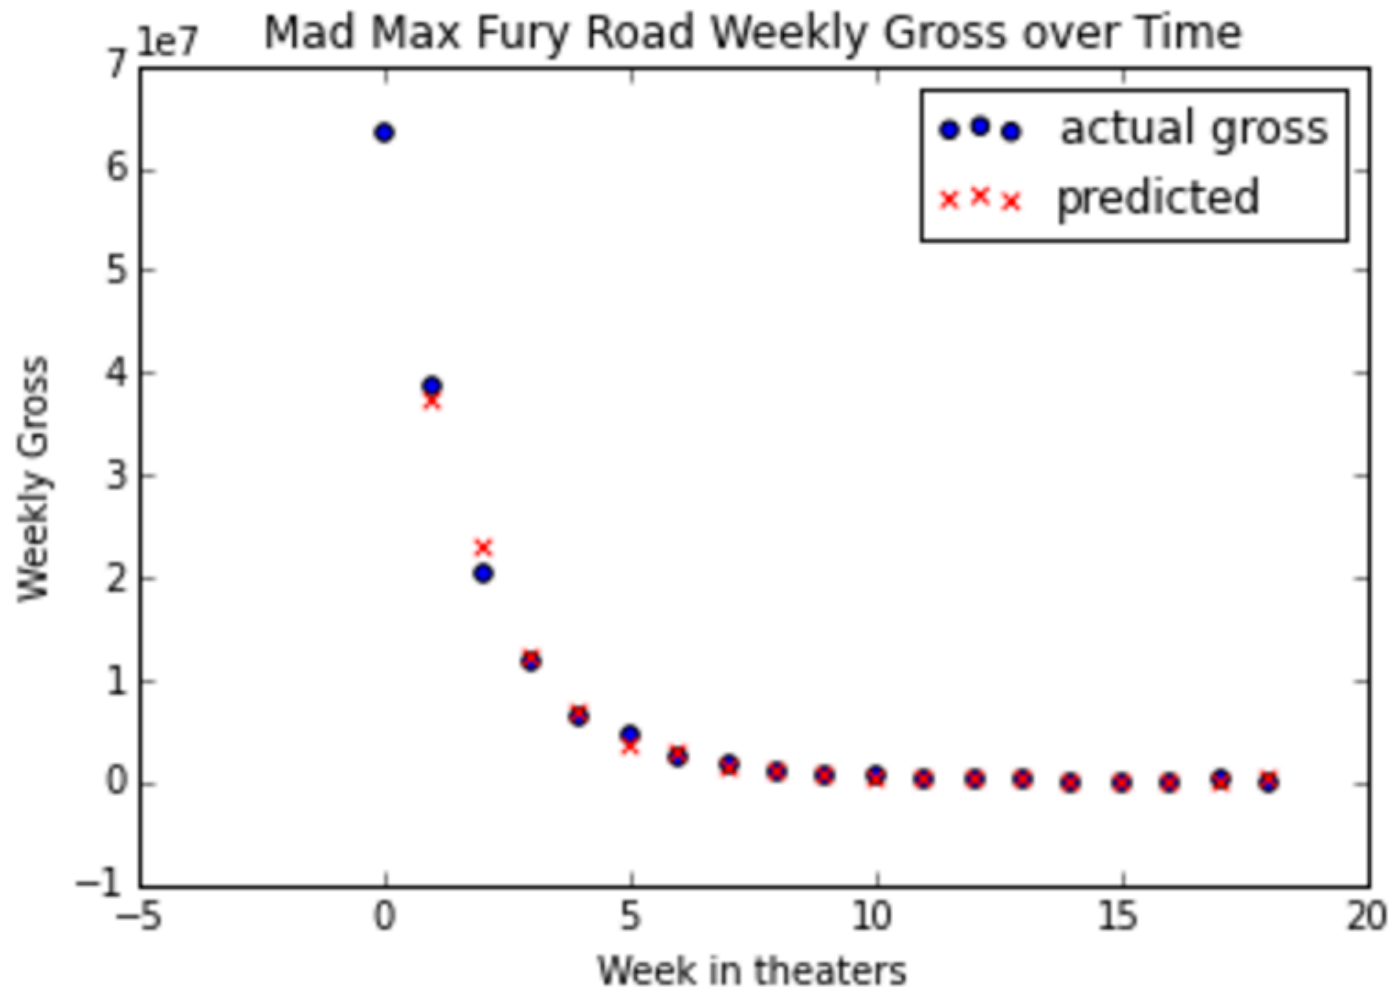
# AR-1

We expect a week's gross for a movie to be  
\$10 Thousand  
+ 60% of what it made last week  
+ luck factor  
(roll the dice for +- \$110 Thousand)

$$x_t = \beta_0 + \beta_1 x_{t-1} + \varepsilon$$


$$\begin{aligned}\beta_0 &= 0.01 \text{million} \\ \beta_1 &= 0.598\end{aligned}$$

# AR-1




# AR-1

Looking at the order of magnitude (log) of the gross instead to be able to see differences in later weeks



# Previously, we were drawing target versus feature

Here is the AR1 model from the same angle: Drawing  
target vs an input feature



# AR-2

What we know  
(features)

What we predict  
(target)

Gross of previous week  
Gross of two weeks ago



Gross of  
this week

Target

Feature 1

Feature 2



	weekly_gross	one_prev_weeks_gross	two_prev_weeks_gross
0	63440279	NaN	NaN
1	38849255	63440279	NaN
2	20544731	38849255	63440279
3	11643562	20544731	38849255
4	6309002	11643562	20544731
5	4555993	6309002	11643562
6	2648047	4555993	6309002
7	1645168	2648047	4555993
8	966275	1645168	2648047
9	601794	966275	1645168
10	663222	601794	966275

AR-2

*AutoRegressive  
model with lag 2*


Two previous points  
as features

# AR-3

and so on and so forth...

What we know  
(features)

What we predict  
(target)



You can do the lag processing yourself, or just use AR models in statsmodels to do it for you

```
import statsmodels.api as sm
```

```
AR1 = sm.tsa.AR(weekly_flu_numbers, freq="W").fit(maxlag=1)
```

```
AR2 = sm.tsa.AR(weekly_flu_numbers, freq="W").fit(maxlag=2)
```

```
AR2 = sm.tsa.AR(weekly_flu_numbers, freq="W").fit(maxlag=3)
```



Big splash in 2008: Google Flu Trends!

**Google can predict flu volume  
by analyzing Google searches!**

Big splash in 2008: Google Flu Trends!

**Google can predict flu volume  
by analyzing Google searches!**

**Majorly Accurate Flu  
Model from Google!**

# Big splash in 2008: Google Flu Trends!

By looking at how many people google for flu related terms, Google can predict flu epidemics!

**Google can predict flu volume  
by analyzing Google searches!**

**Majorly Accurate Flu  
Model from Google!**

# Big splash in 2008: Google Flu Trends!


Historical estimates

See data for: United States

## United States Flu Activity

Influenza estimate

● Google Flu Trends estimate ● United States data



United States: Influenza-like illness (ILI) data provided publicly by the [U.S. Centers for Disease Control](#).

Neat!

OMG AMAZANNGG!  
It's, like, the future!

Historical estimates


See data for:

United States

## United States Flu Activity

Influenza estimate

● Google Flu Trends estimate ● United States data



United States: Influenza-like illness (ILI) data provided publicly by the [U.S. Centers for Disease Control](#).

# Google Flu Model

What we know  
(features)

Search volume  
for flu related  
terms this week



What we predict  
(target)

Number of  
flu cases  
this week

The bubble burst when we realized that  
an autoregressive model can produce the  
same accuracy

What we know  
(features)

Number of  
flu cases  
last week



What we predict  
(target)

Number of  
flu cases  
this week

And actually adding google search data as a new feature gives only a tiny improvement in performance

What we know  
(features)

Search volume for  
flu related terms  
this week

Number of flu  
cases last week



What we predict  
(target)

Number of  
flu cases  
this week



# AR can be powerful

What we know  
(features)

Number of  
flu cases  
last week



What we predict  
(target)

Number of  
flu cases  
this week

This has been  
available even  
before the  
Internet

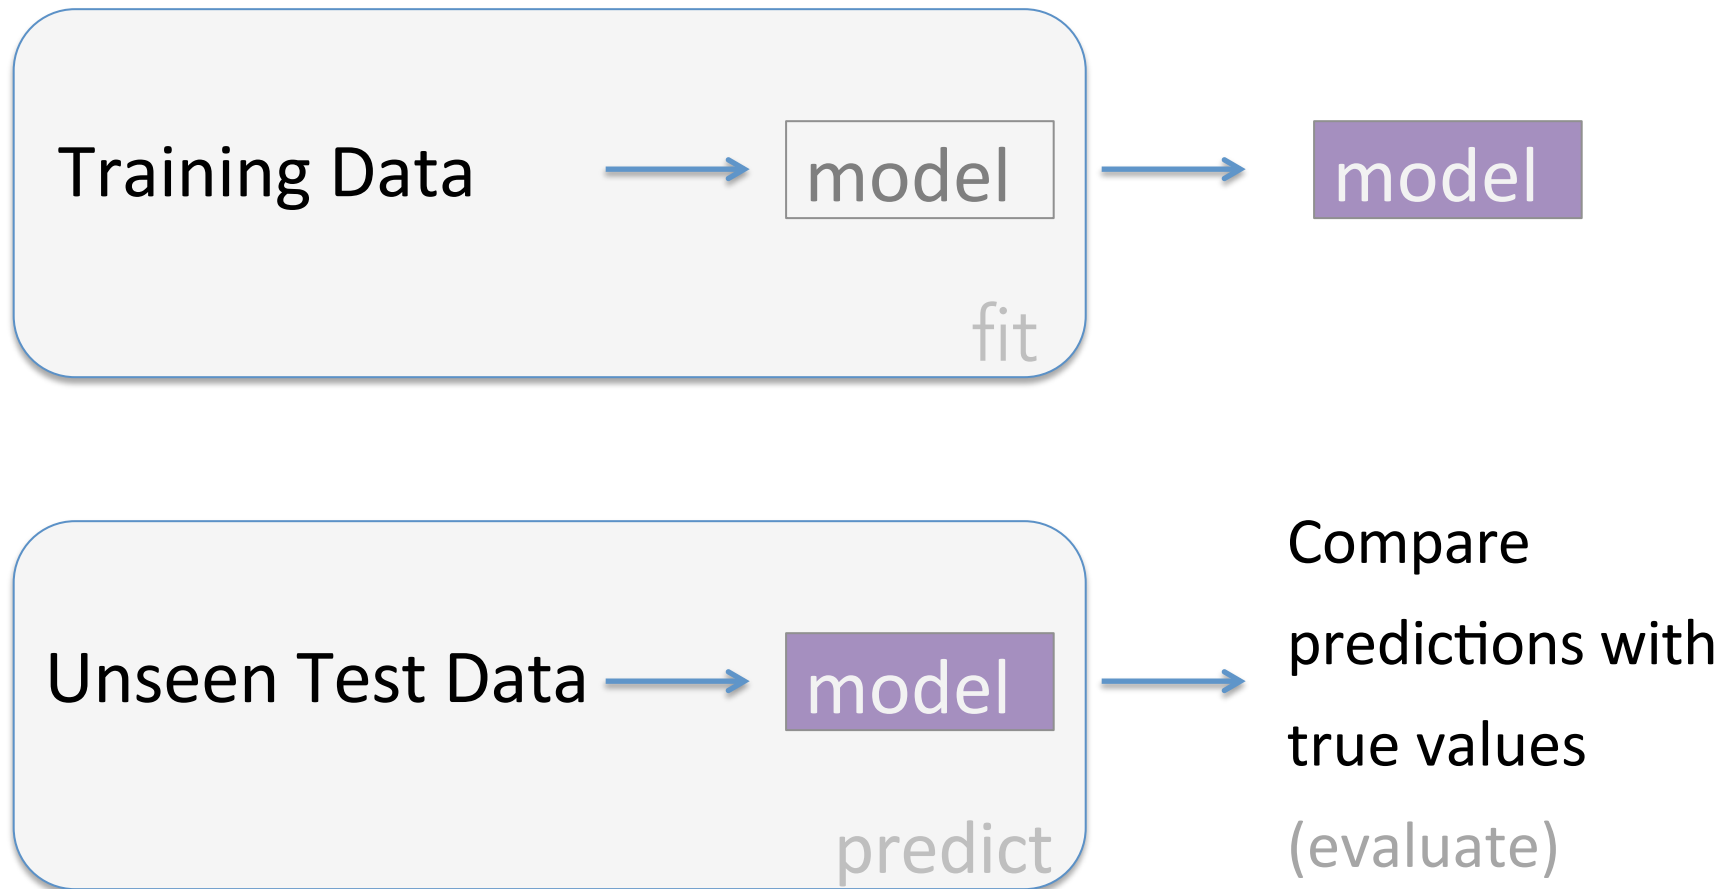
Think of lagged-features just as another potential source of usable information in a standard regression problem.

You can combine them with other features as well.


AR models can be quite powerful for **short term forecasting** (next time point), but their efficiency goes down when you **try to forecast long term** (as errors propagate from each next point you predict)

# How do we do the train/test split?

As with any regression problem, it is important to measure performance on a separate test set.




In previous cases, we could just randomly pick some points and set them aside as a test set



66% Training set

33% Test set


In previous cases, we could just randomly pick some points and set them aside as a test set



66% Training set  
33% Test set


# What about time series?

These points aren't independent




# Depends on your use case!

Test set should simulate exactly what type of points you will make predictions on with the model



# Train/test split for time series


**Case 1:** We model a with past observations to make predictions on how it will continue.





# Train/test split for time series

**Case 1:** We model a with past observations to make predictions on how it will continue.




Mad Max has been running for 18 weeks.

Each week, I want to ask the model to forecast next week's gross, to decide if it's worth keeping it on one more week


# Train/test split for time series

**Case 1:** We model a with past observations to make predictions on how it will continue.



# Train/test split for time series


**Case 1:** We model a with past observations to make predictions on how it will continue.



Predict the future

# Train/test split for time series


**Case 1:** We model a with past observations to make predictions on how it will continue.



Your training/test split should simulate building a model on past values and testing them on newer values!

# Train/test split for time series


**Case 1:** We model a with past observations to make predictions on how it will continue.



You can even try it on different windows (kind of similar to cross-validation, but not quite)

# Train/test split for time series


**Case 1:** We model a with past observations to make predictions on how it will continue.



You can even try it on different windows (kind of similar to cross-validation, but not quite)

# Train/test split for time series

**Case 1:** We model a with past observations to make predictions on how it will continue.



You can even try it on different windows (kind of similar to cross-validation, but not quite)

# Train/test split for time series

**Case 2:** We model gross decay over time in movies in the past to predict how it'll go in new movies.



?

We just opened The Force Awakens. We need a model that can study all past movies and predict what gross to expect each week.



# Train/test split for time series

**Case 2:** We model gross decay over time in movies in the past to predict how it'll go in new movies.



Model  
observed movies

?



We just opened The Force Awakens. We need a model that can study all past movies and predict what gross to expect each week.

Predict  
new movie

# Train/test split for time series

**Case 2:** We model gross decay over time in movies in the past to predict how it'll go in new movies.



Randomly assign movies into training and test sets, fit AR model to the points from the movies in the training set, evaluate performance on test set movies

# Train/test split for time series

**Depending on your use case, set aside a test set that simulates the way in which you will ask for predictions**

# How to choose the order of AR

AR1? AR2? AR3? Which do I use?

Previous point's value, two-previous point's value, etc.  
are different features.

Just like in any other regression, **you will try different feature sets and choose the model with the test performance.**

Using more lag-features can improve performance, too many can lead to overfitting. Depends on the specific problem.

# How to choose the order of AR

AR1? AR2? AR3? Which do I use?

Note:

You can also look at the autocorrelation function to inform your process of choosing AR order, but in the end, a hypothesis-driven trial-and-error process with performance evaluation on separate test sets will be the ultimate best way of determining order.

We will touch upon autocorrelation functions later in the bootcamp, when we revisit Time Series.


# Autocorrelation

Another input to understand the time series better and help choose the order of the AR model

# Autocorrelation

Imagine a simple regression model. No time series.


We need to predict total movie gross, and we have several features we can use. We are trying to understand the system and choose features.



# Autocorrelation


A good approach is to plot each feature by itself versus the target, and look at their correlation. This way we can see how much information each feature carries.

Movie  
Gross




Budget

Movie  
Gross




Runtime

Movie  
Gross



RT score

Movie  
Gross


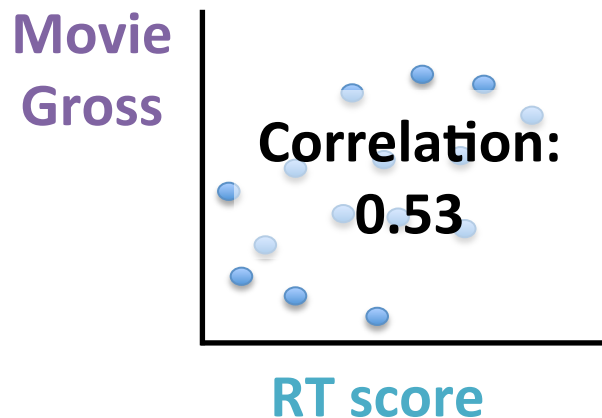




Number of words in the title







# Autocorrelation

A good approach is to plot each feature by itself versus the target, and look at their correlation. This way we can see how much information each feature carries.




# Autocorrelation

Of course, this by itself cannot be the sole determining information in selecting the feature set, but it helps us understand the contribution of information!



# Autocorrelation


Now imagine our weekly gross problem again.  
Basically, trying to figure out what order AR to use is the same as trying to figure out which features to use.



# Autocorrelation


It is still a good idea to understand the correlations between each of these features and the target.

This  
Week's  
Gross




Gross from last week

This  
Week's  
Gross




Gross from 3 weeks ago

This  
Week's  
Gross



Gross from 2 weeks ago





This  
Week's  
Gross



Gross from 4 weeks ago

# Autocorrelation

It is still a good idea to understand the correlations between each of these features and the target.



# Autocorrelation

The autocorrelation function basically condenses this information into a single plot

**Correlation:**  
**0.69**


**Correlation:**  
**0.31**

**Correlation:**  
**0.48**

**Correlation:**  
**0.10**


# Autocorrelation

The autocorrelation function basically condenses this information into a single plot



# Autocorrelation

The autocorrelation function basically condenses this information into a single plot





# Autocorrelation

```
import statsmodels as sm  
sm.tsa.stattools.acf(timeseries)
```


creates the autocorrelation function for you.

# Preprocessing: Moving Average


specific to time-series

# Preprocessing: Moving Average

specific to time-series




On a day to day basis, weight fluctuates up and down.  
On a month to month basis, it is going down.



On a day to day basis, weight fluctuates up and down.


On a month to month basis, it is going down.

If we can make solid assumptions about what is noise  
and what is a trend, we can make the modeling  
problem easier



## Assumption: Day to day fluctuations are “noise”.


What that really means that we don't aim to build a model that can predict such daily changes. We want to model longer-time-scale trends.



**Important:** We are not making a claim about what we think is “really” happening underneath.


We are declaring what we choose to model.

As long as you don’t forget these decisions when you use the model for predictions, and stay true to them, they are fine.



**Assumption: Day to day fluctuations are “noise”.**


With this assumption I can remove some of that noise with a moving average window.






**Assumption: Day to day fluctuations are “noise”.**

With this assumption I can remove some of that noise with a moving average window.




**Assumption: Day to day fluctuations are “noise”.**

With this assumption I can remove some of that noise with a moving average window.




**Assumption: Day to day fluctuations are “noise”.**

With this assumption I can remove some of that noise with a moving average window.




**Assumption: Day to day fluctuations are “noise”.**

With this assumption I can remove some of that noise with a moving average window.




**Assumption: Day to day fluctuations are “noise”.**

With this assumption I can remove some of that noise with a moving average window.




**Assumption: Day to day fluctuations are “noise”.**

With this assumption I can remove some of that noise with a moving average window.




**Assumption: Day to day fluctuations are “noise”.**

With this assumption I can remove some of that noise with a moving average window.



# Exponential Moving Average


Instead of weighing each point in the window the same amount, average all past points, but decrease the weight exponentially as we go further back.





# Preprocessing: Moving Average

Preprocessing can help avoid overfitting to noise.  
Always be mindful of the assumptions you made.



# **Warning!**

## Moving Average Smoothing is not Moving Average Models (MA)

What we covered is a preprocessing technique in modeling time series.

Moving Average Models are (just like AR models) a framework to model time series.

They are inaccurately named (they don't involve averages). We will cover them briefly in the future.